



When ECM Meets the Semantic Web

20 Oct 2011 - Olivier Grisel & Stefane Fermigier

Business Motivations

2

NUXEO WORLD PARIS 2011



Information society

From Wikipedia, the free encyclopedia

For other uses, see [Information society \(disambiguation\)](#).

The aim of the **information society** is to gain competitive advantage internationally through using IT in a creative and productive way. An **information society** is a [society](#) in which the creation, distribution, diffusion, use, integration and manipulation of [information](#) is a significant economic, political, and cultural activity. The [knowledge economy](#) is its economic counterpart whereby wealth is created through the economic exploitation of understanding. People that have the means to partake in this form of society are sometimes called [digital citizens](#). As Beniger^[1] shows, this is one of many dozen labels that have been identified to suggest that we are entering a new phase of society.

Source:Wikipedia

Knowledge economy

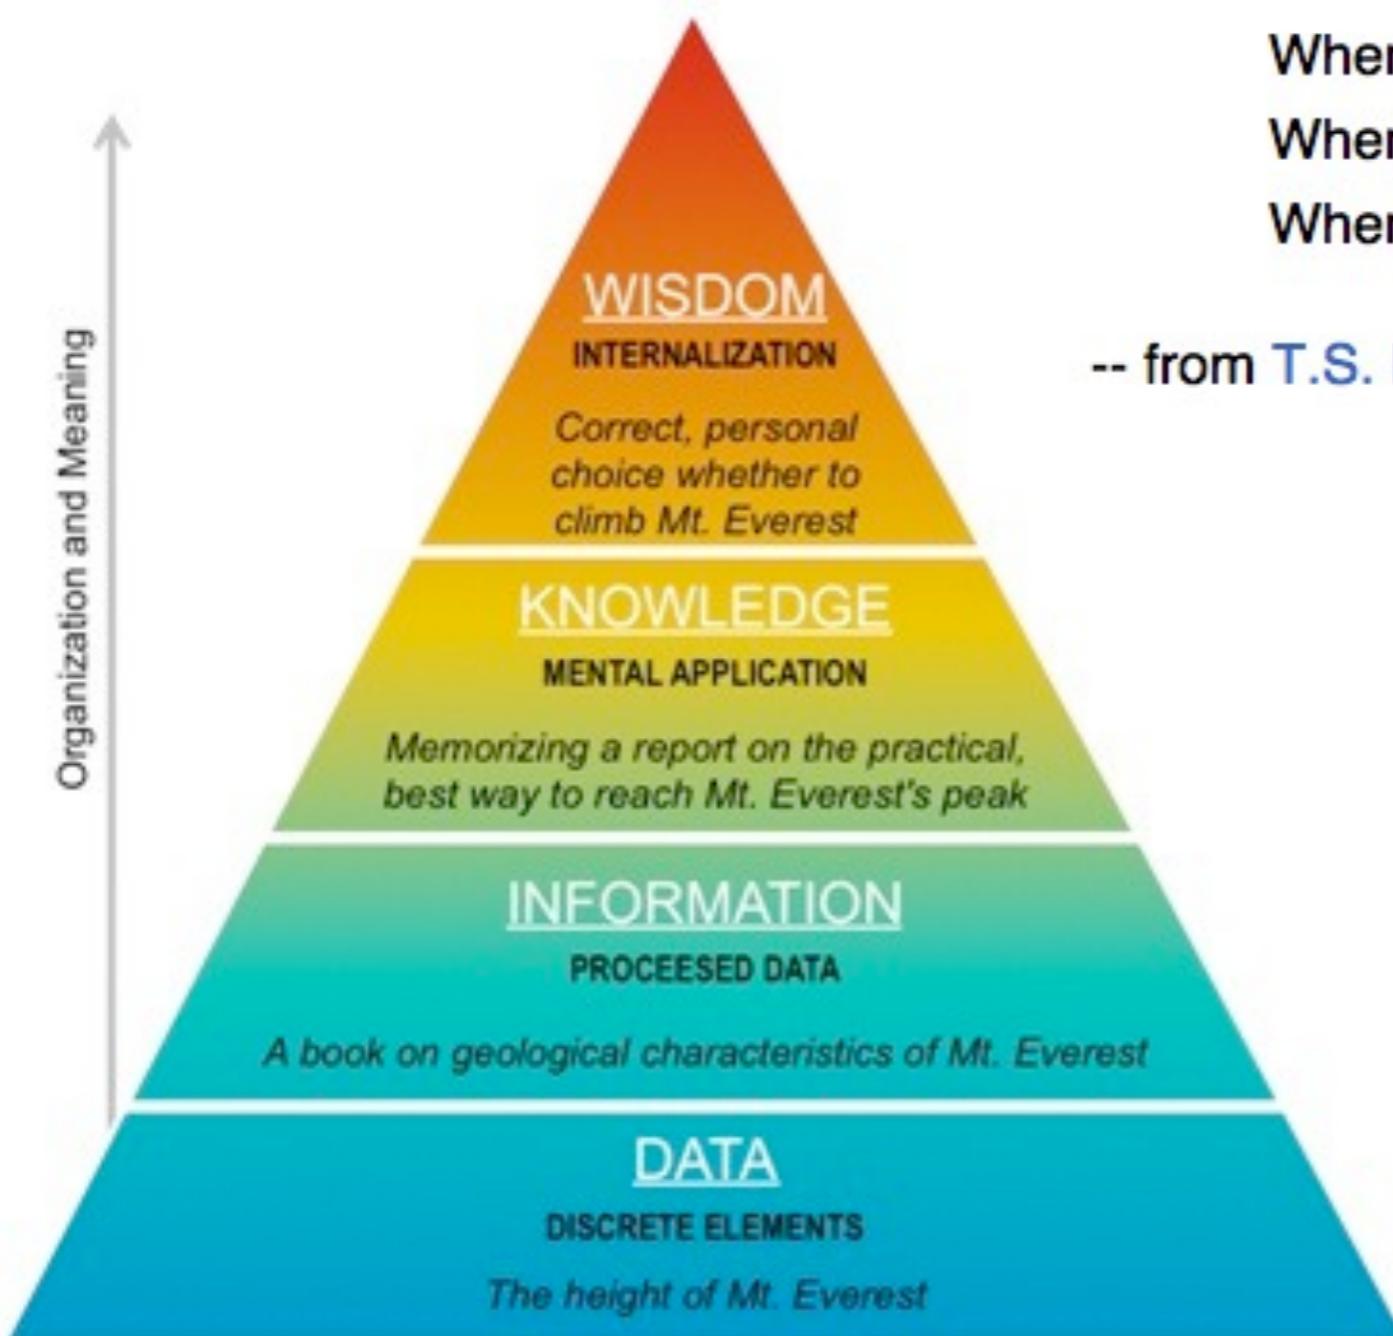
From Wikipedia, the free encyclopedia

The **knowledge economy** is a term that refers either to an **economy of knowledge** focused on the production and management of knowledge in the frame of **economic** constraints, or to a **knowledge-based economy**. In the second meaning, more frequently used, it refers to the use of **knowledge** technologies (such as **knowledge engineering** and **knowledge management**) to produce **economic** benefits as well as job creation. The phrase was popularized by **Peter Drucker** as the title of Chapter 12 in his book *The Age of Discontinuity*, And, with a footnote in the text, Drucker attributes the phrase to economist **Fritz Machlup**.^[1]

The essential difference is that in a *knowledge economy*, knowledge is a product, while in a *knowledge-based economy*, knowledge is a tool. This difference is not yet well distinguished in

Source:Wikipedia

The DIKW hierarchy



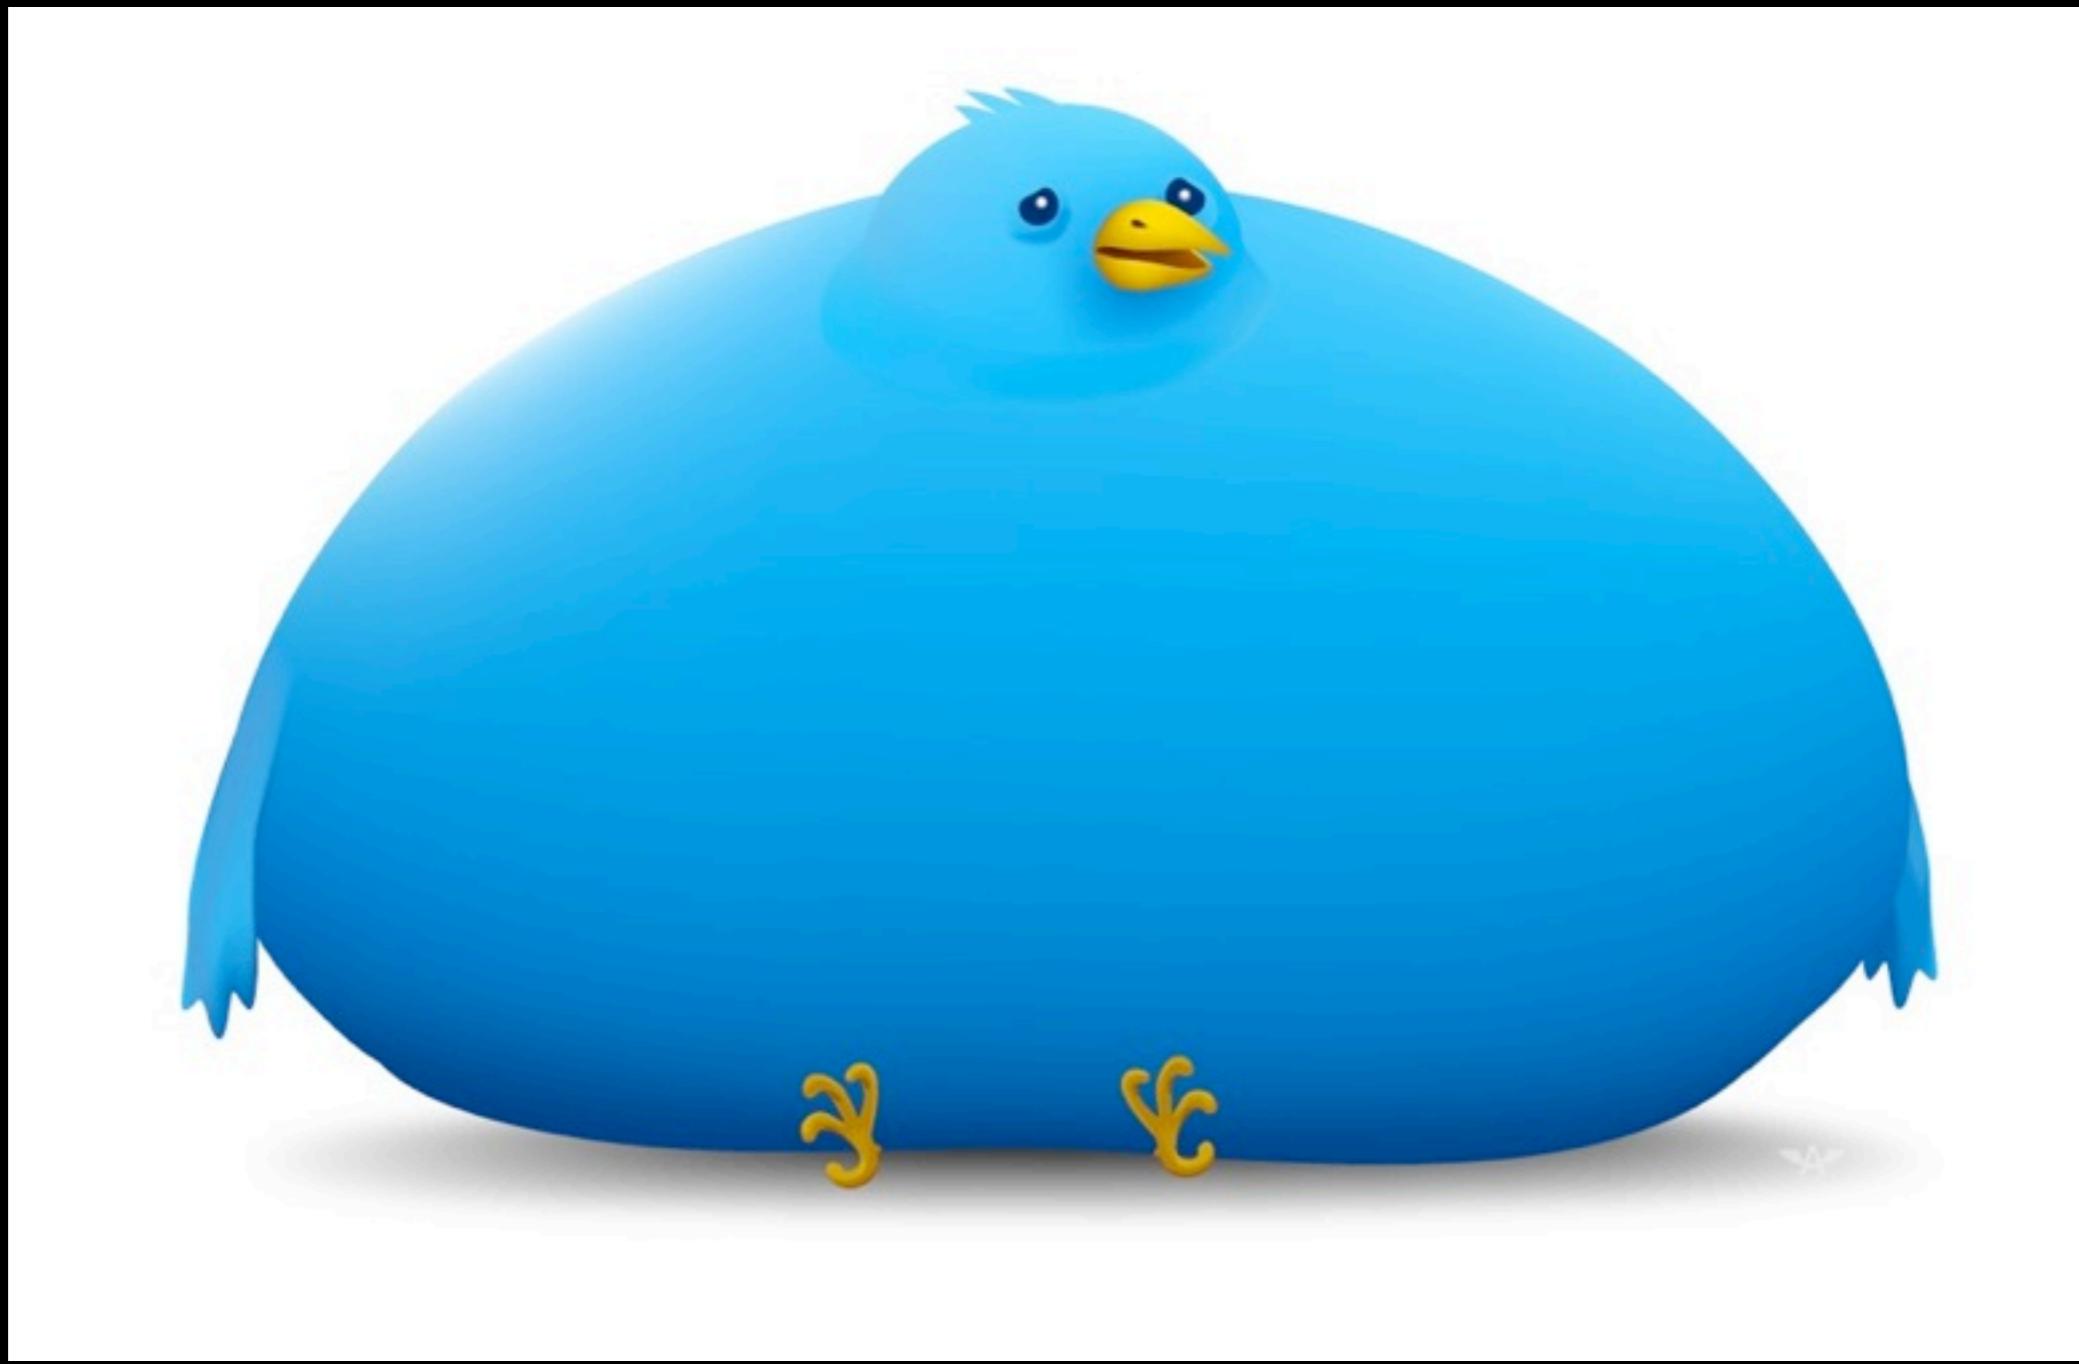
Where is the Life we have lost in living?
Where is the wisdom we have lost in knowledge?
Where is the knowledge we have lost in information?

-- from T.S. Eliot, "Choruses from 'The Rock'"



But every coin has
another side

Infobesity!



A few figures

- 50% more data / content / information produced every year
- 1.8 zettabytes of data produced in 2011 (= 1 billion terabytes)
- Employees are drowning in a sea of email, status messages, etc., and spend on average more than 6 hours / weeks unsuccessfully searching for or recreating lost documents

A Solution: the Semantic Web

9

NUXEO WORLD PARIS 2011



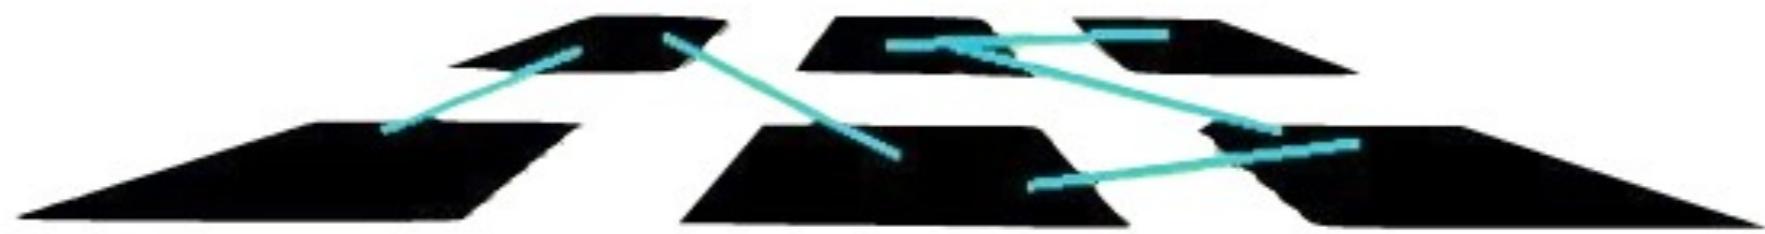
A Brief History of the Web

- **Web 1.0** (1990-now): web of sites and pages, aka the *World Wide Web*
- **Web 2.0** (2000-now): web of people and of participation, aka the *Social Web* (Blogs, RSS, tags, Facebook, Wikipedia, etc.)
- **Web 3.0** (2010-now): web of data, of meaning and connected knowledge, aka the *Semantic Web*





11



“To a computer, then, the web is a **flat, boring** world devoid of **meaning**”

Tim Berners Lee, <http://www.w3.org/Talks/WWW94Tim/>

12

NUXEO WORLD PARIS 2011

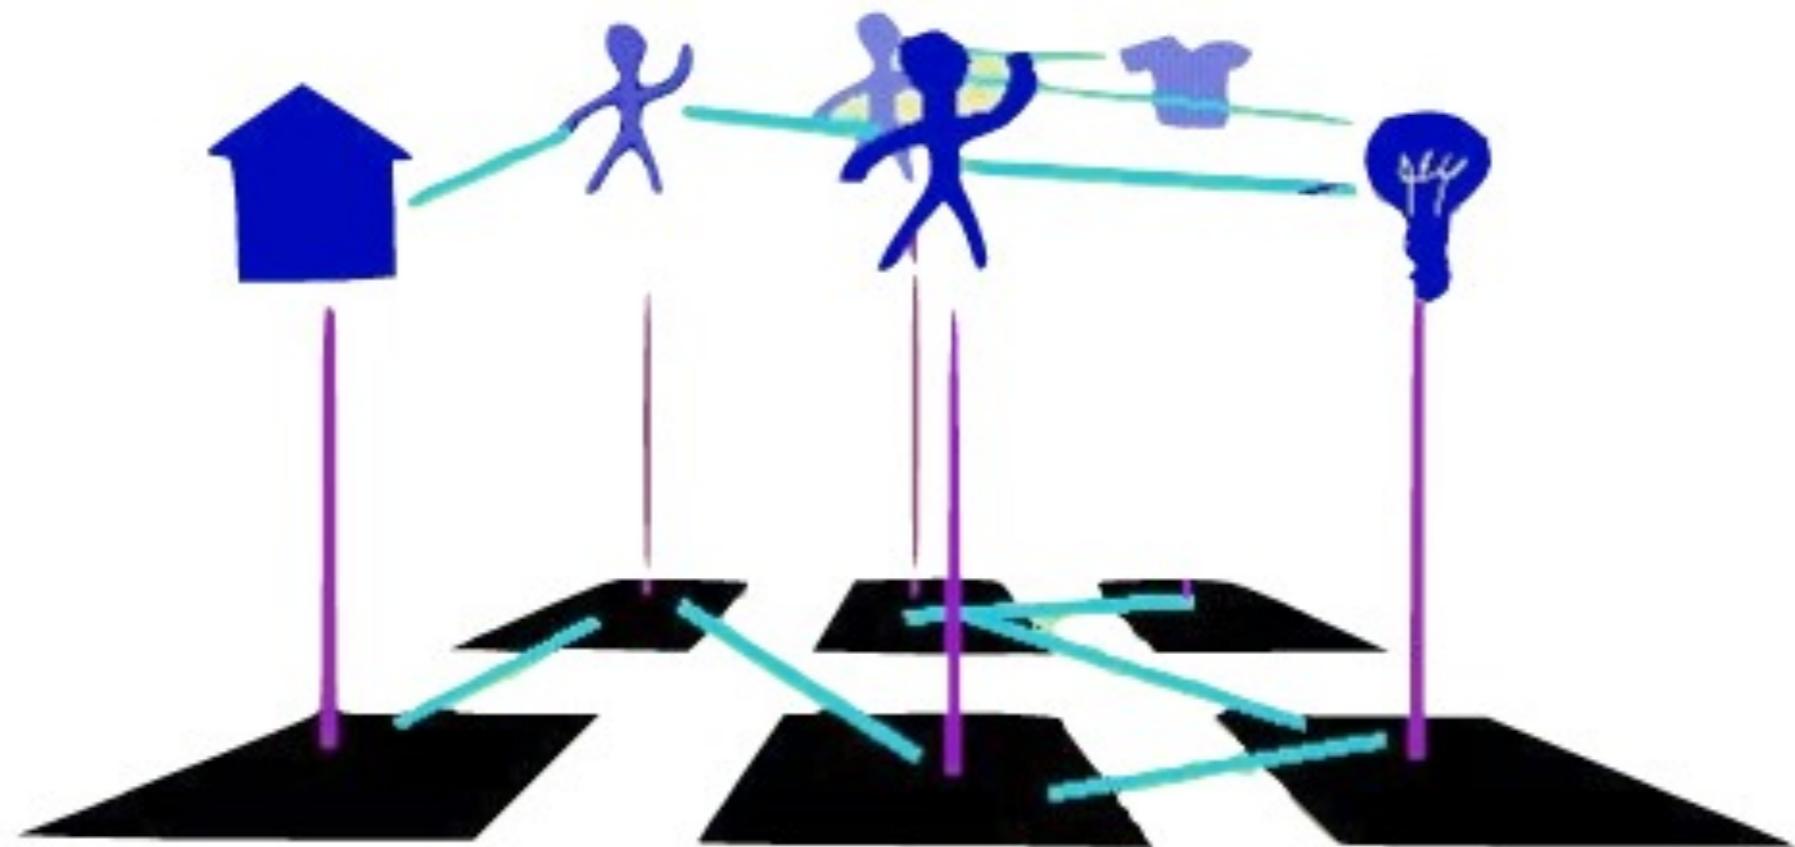




“This is a pity, as in fact **documents** on the web describe **real objects** and imaginary **concepts**, and give particular **relationships** between them”

Tim Berners Lee, <http://www.w3.org/Talks/WWW94Tim/>

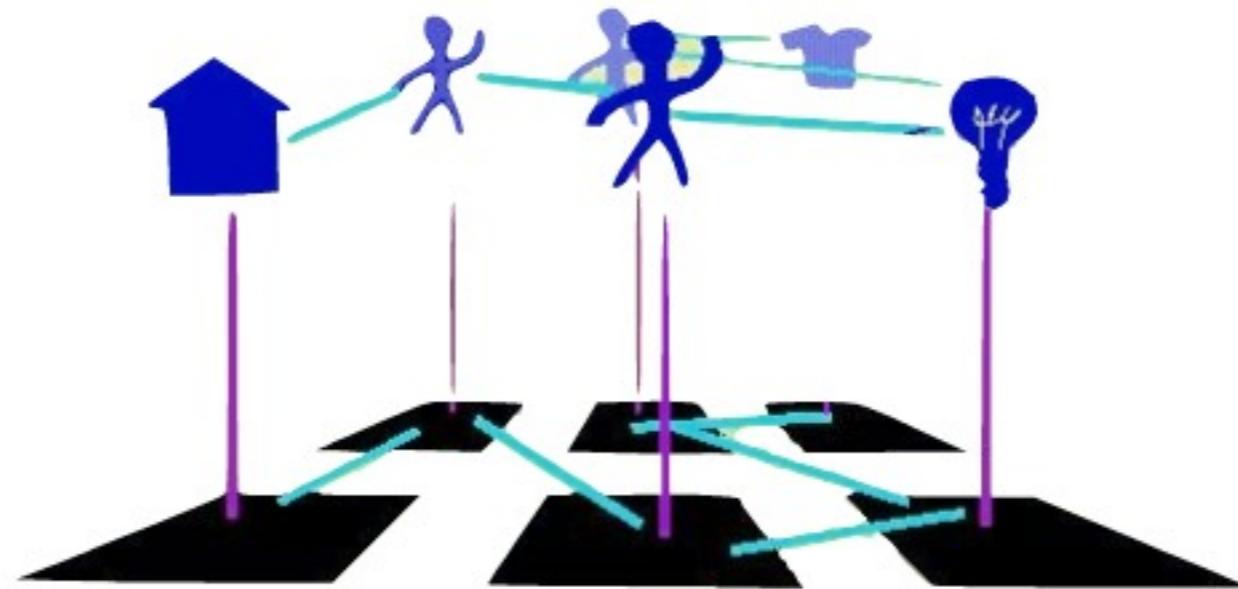




“Adding semantics to the web involves two things: allowing **documents** which have information in **machine-readable** forms, and allowing **links** to be created with **relationship values**.”

Tim Berners Lee, <http://www.w3.org/Talks/WWW94Tim/>





“The Semantic Web is not a separate Web but an **extension** of the current one, in which information is given well-defined **meaning**, better enabling **computers and people** to work in cooperation.”

Tim Berners Lee, <http://www.w3.org/Talks/WWW94Tim/>



Means and Tools

16

NUXEO WORLD PARIS 2011



4 stages

- **Extract** meaning from raw data / content
- **Connect** information to form knowledge
- **Reason** about this knowledge
- **Present** this knowledge in actionable form



Extracting

- Leverage metadata embedded in or associated with documents (when they exist)
- Or use machine learning, NLP (Natural Language Processing) and image processing algorithms to **extract meaning** from text / images
- Examples include: named entities extraction, automatic categorization / tagging, sentiment analysis, etc.



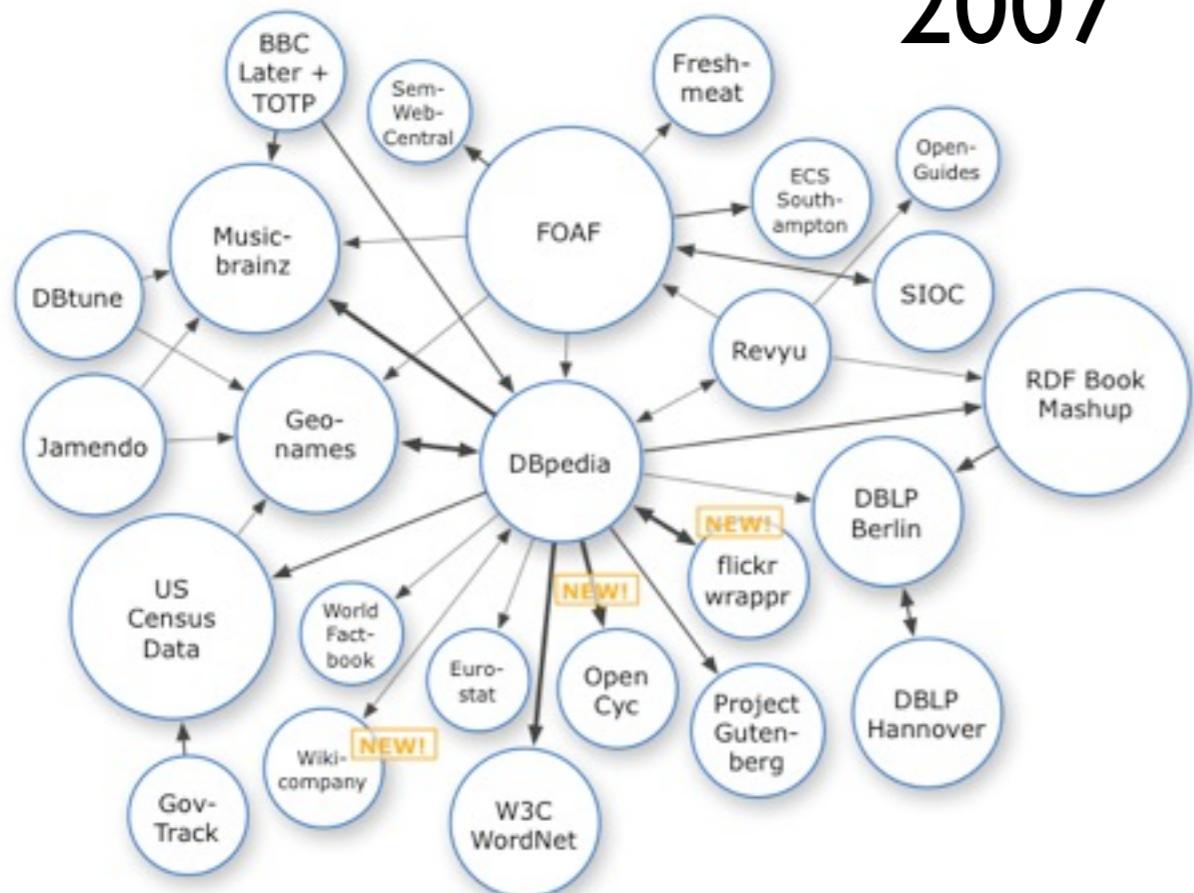
Interlude: Linked Open Data

19

NUXEO WORLD PARIS 2011



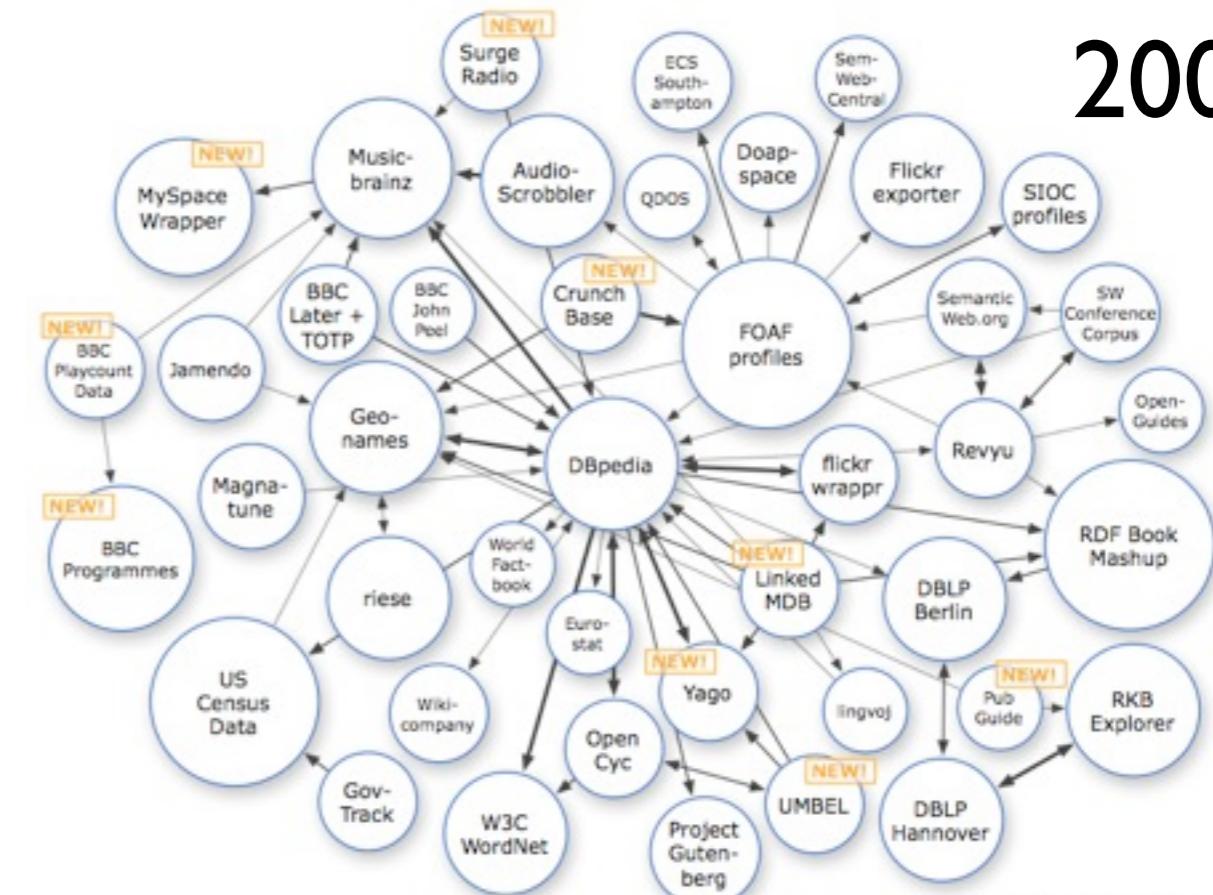
2007



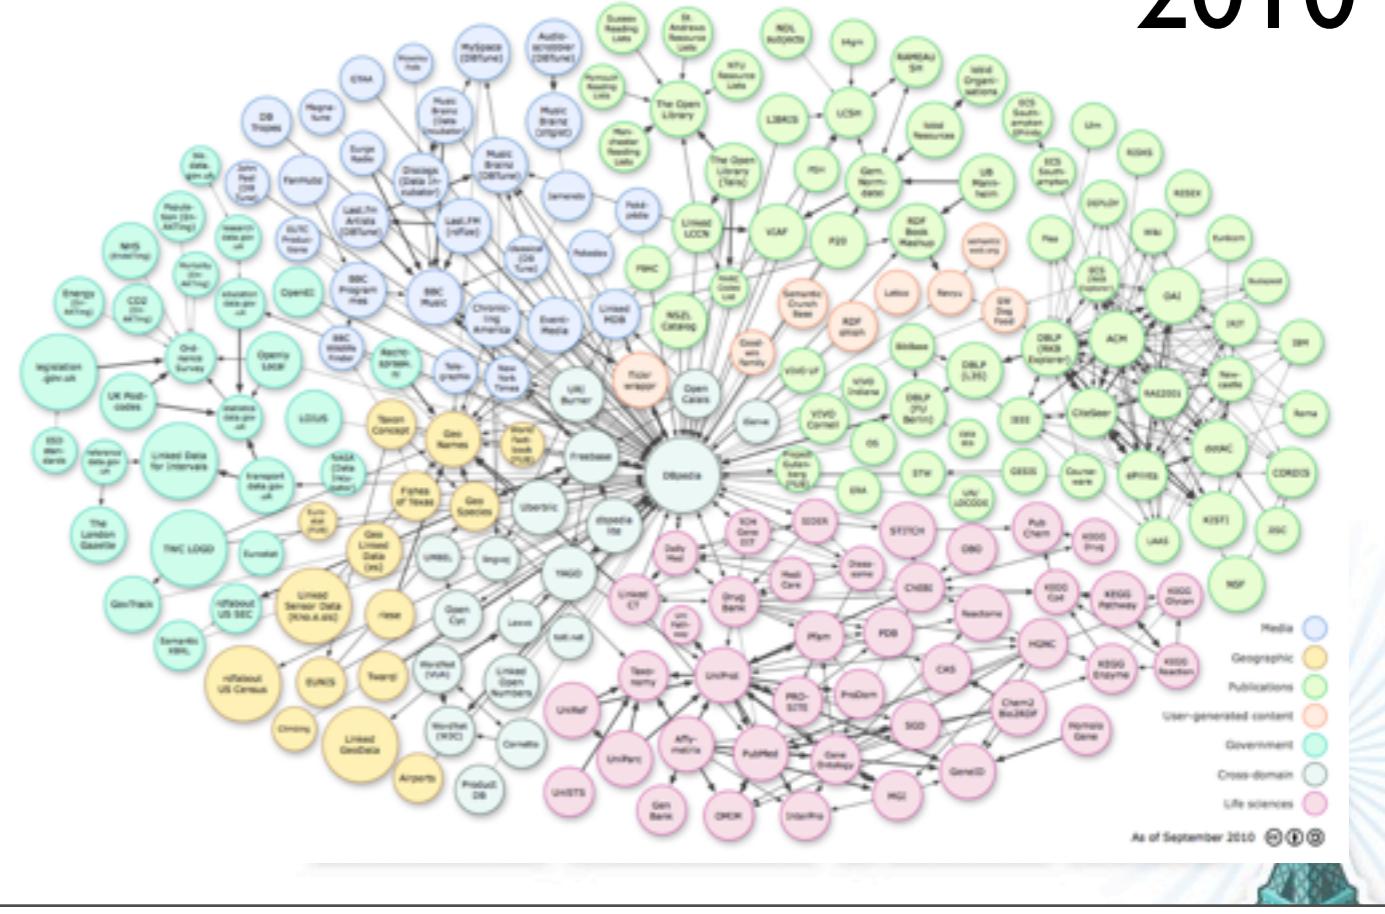
2009



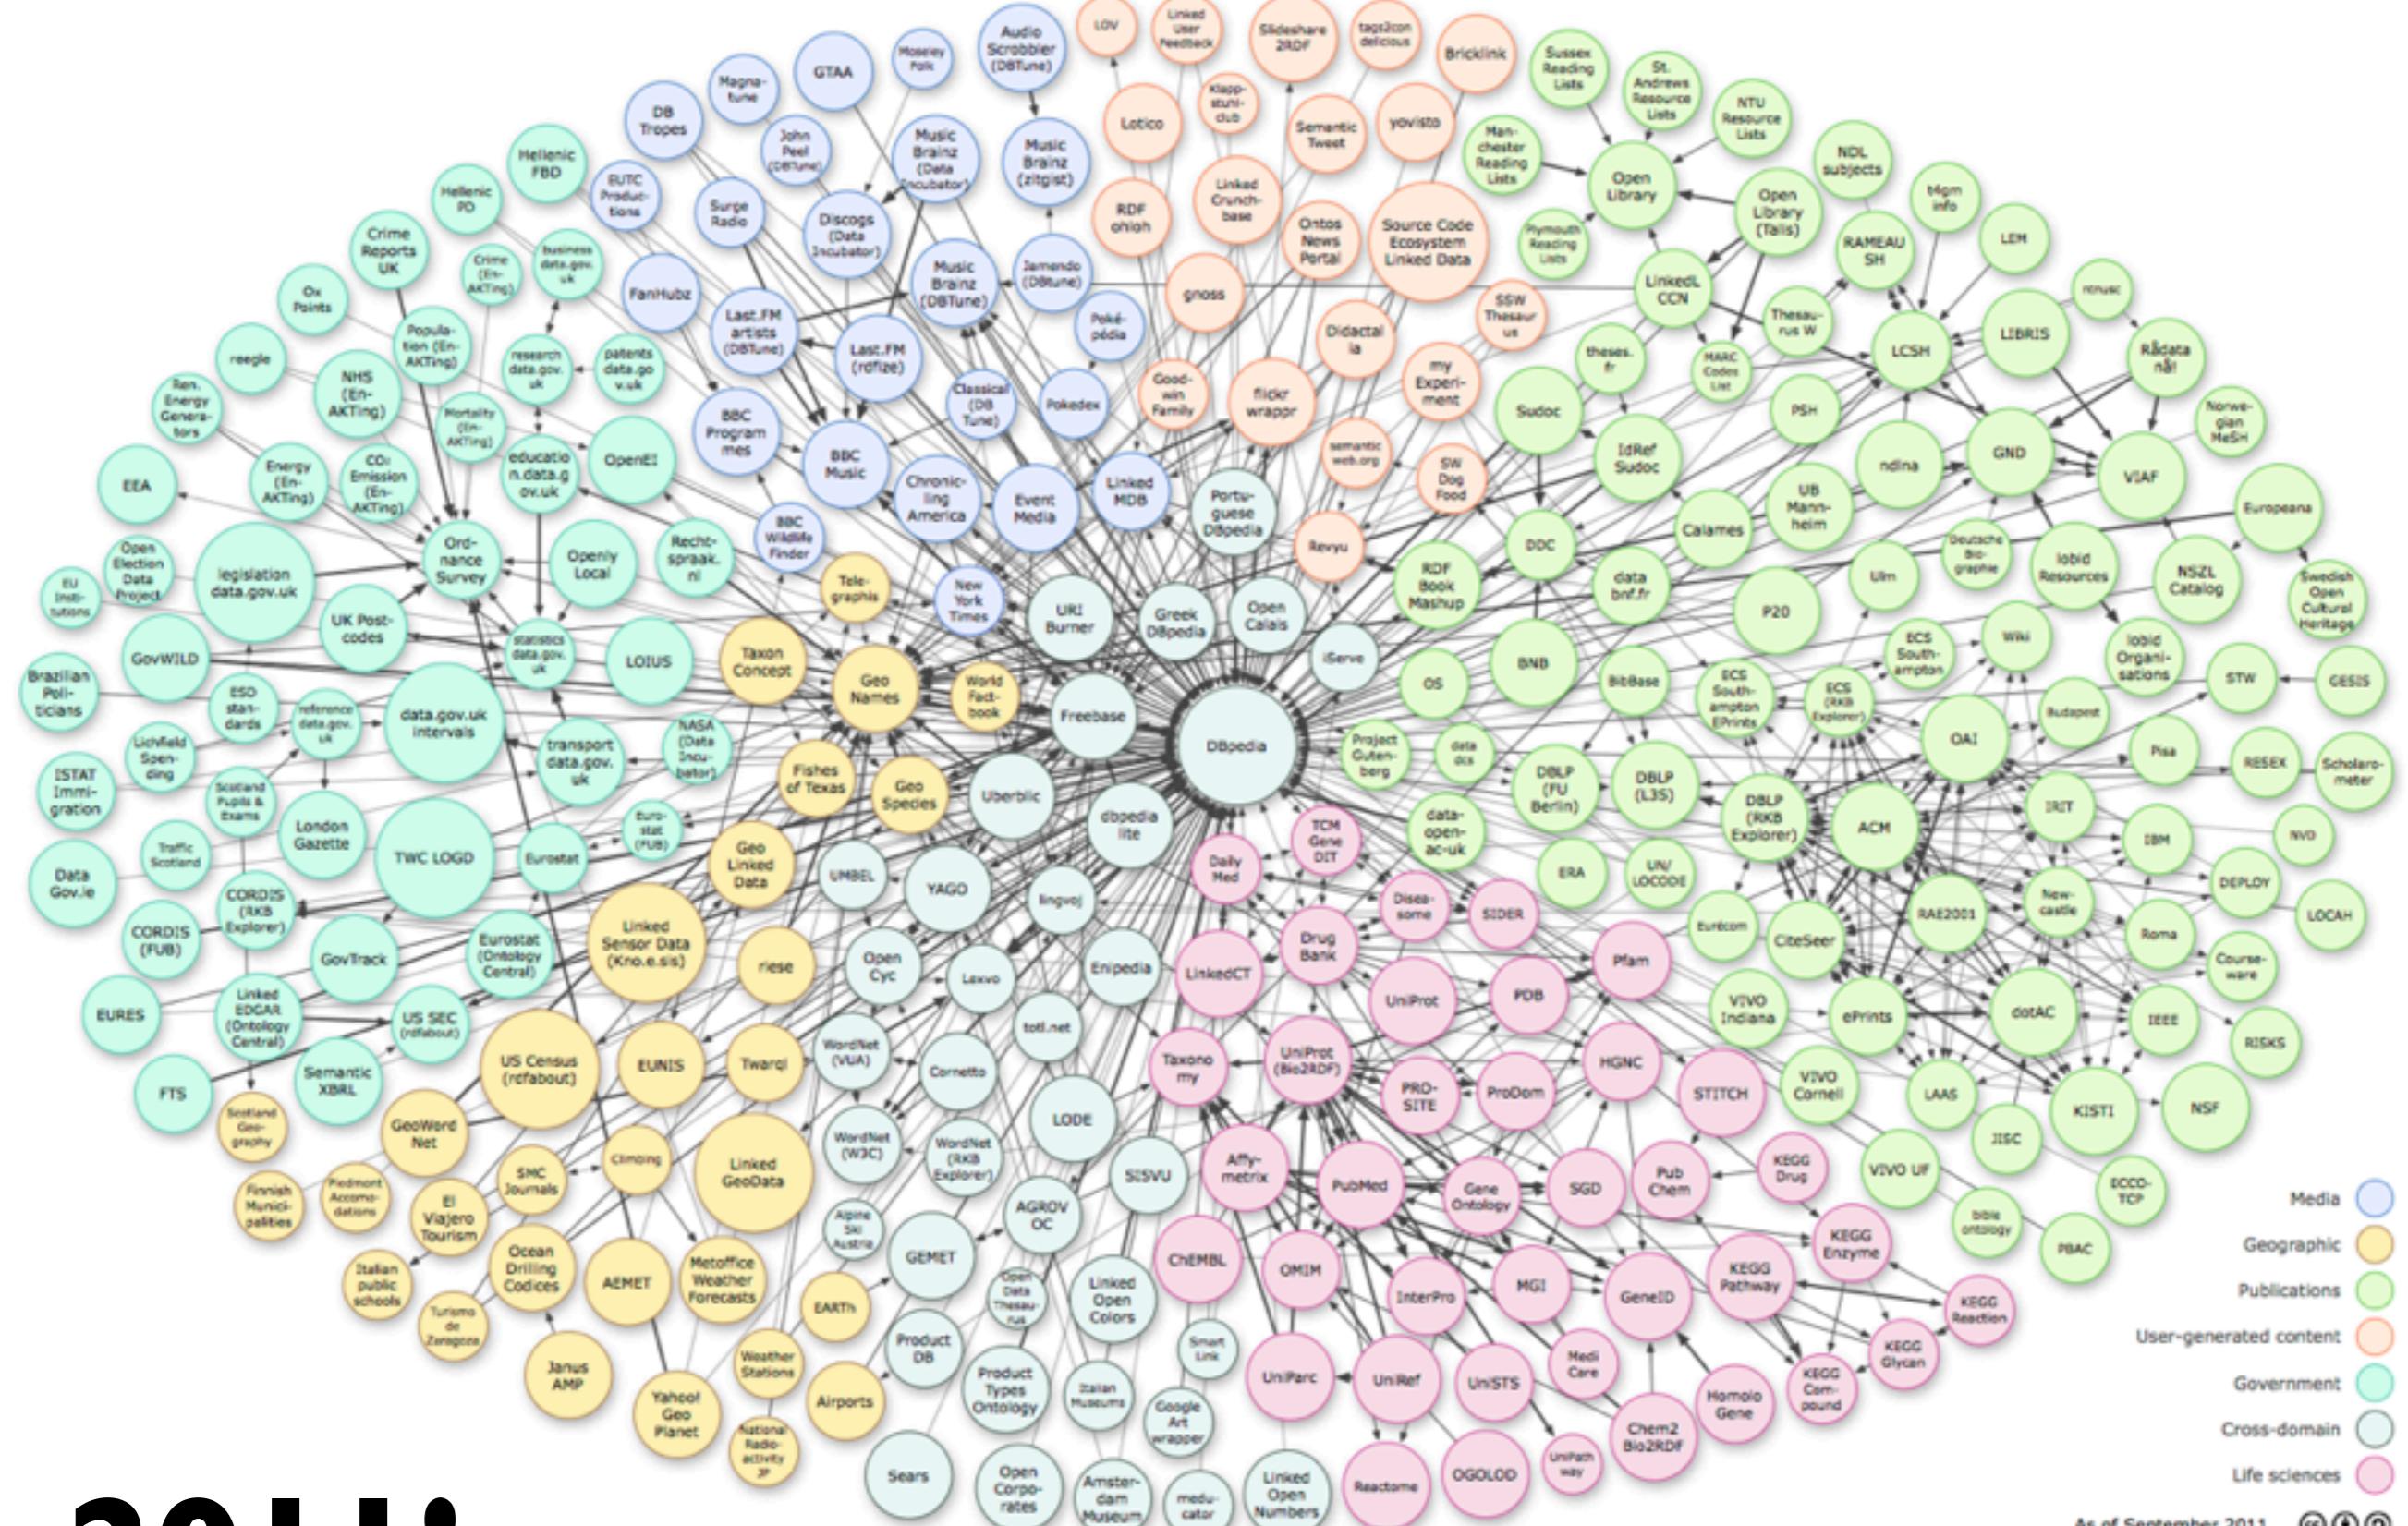
As of July 2009



As of September 2008



Page 1 of 1



2011!

As of September 2011



NUXEO WORLD PARIS 2011



Linking

- Many Linked Open Data repositories have been made available over the last 10 years
- RDF and graph database systems are now available to manage this huge mass of information (billions of triples)
- **Match** information extracted from content with these public (or internal) data/knowledge bases



Reasoning

- When you are working on reliable metadata (ex: RDFa embedded in web pages), you can use rule / inference engines to infer **actionable knowledge** from your content (ex: shopping recommendation engine)
- Rules can also be used to **clean up** / flag errors when working with unreliable (e.g. automatically extracted) information



Presenting

- Allow the users of your system to interact with the knowledge thus extracted or produced, in a way that allows them to do their jobs better
- A smart presentation system solves the information overload issue by **contextualizing** the information, i.e. presenting only information relevant to what the user is currently doing



R&D Projects Involving Nuxeo

25

NUXEO WORLD PARIS 2011



IKS project



- European R&D project under the FP7, with 13 partners (6 SMEs) and a 8.5M EUR budget
- Goal: create a semantic software “stack” that will be used by CMS vendors to add semantic features to their products
- Started in Jan. 2009, will last until Dec. 2012
- First tangible result: **Apache Stanbol** (more about this later)



SAMAR project



- French collaborative R&D project with 10 partners, and a 4.5M EUR budget
- Goal: create a platform for managing multimedia content in arabic, for news agencies such as AFP
- Will include: automated translation, named entities extraction, content classification
- First results: integration between **Nuxeo** and **Temis** (more later)



State of the Art Semantic ECM at Nuxeo

28

NUXEO WORLD PARIS 2011



The Semantic Engine

- From unstructured content to Knowledge
- Language guessing
- Topic classification (Business, Sports, Media, ...)
- Named Entities extraction and linking
- Relationships and properties extraction



Demo time!



31

NUXEO WORLD PARIS 2011





Web View REST API

Enhancement Engines

There are currently 3 active engines.

- ✿ org.apache.stanbol.enhancer.engines.opennlp.impl.NamedEntityExtractionEnhancementEngine
- ✿ org.apache.stanbol.enhancer.engines.autotagging.impl.EntityMentionEnhancementEngine
- ✿ org.apache.stanbol.enhancer.jersey.cache.CachingDereferencerEngine

You can enable, disable and deploy new engines using the [OSGi console](#).

Paste some text below and submit the form to let the active engines enhance it:

John Smith works at Smith Consulting in London, United Kingdom.

Output format: [JSON-LD](#) [Run engines](#)

Extracted entities

People



John Smith

Organizations



Smith Consulting

Places



London



United Kingdom



John Smith



Smith Consulting



London

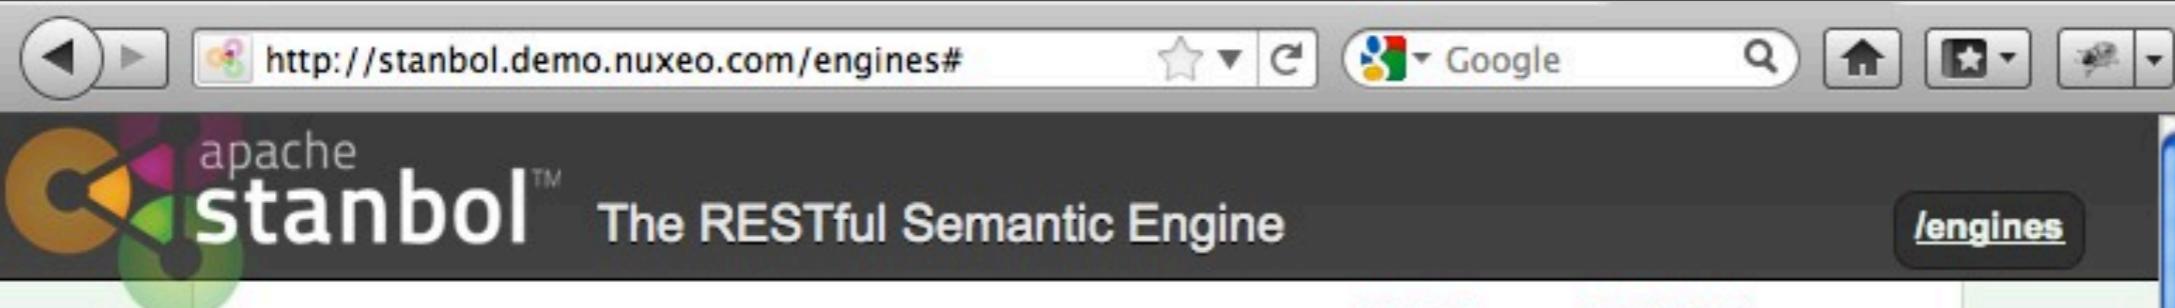


United Kingdom



Raw RDF output

```
{  
  "@": [  
    {  
      "@": "<urn:enhancement-4cf4870a-fc4e-4418-c966-ae4f0a873854>",  
      "a": [
```



The RESTful Semantic Engine

Web View REST API /engines

Enhancement Engines

Stateless REST analysis

This stateless interface allows the caller to submit content to the Stanbol enhancer engines and get the resulting enhancements formatted as RDF at once without storing anything on the server-side.

The content to analyze should be sent in a POST request with the mimetype specified in the Content-type header. The response will hold the RDF enhancement serialized in the format specified in the Accept header:

```
curl -X POST -H "Accept: text/turtle" -H "Content-type: text/plain" \
      --data "John Smith was born in London." http://stanbol.demo.nuxeo.com
/engines
```

The list of mimetypes accepted as inputs depends on the deployed engines. By default only text/plain content will be analyzed

Stanbol enhancer is able to serialize the response in the following RDF formats:

- application/json (JSON-LD)
- application/rdf+xml (RDF/XML)
- application/rdf+json (RDF/JSON)
- text/turtle (Turtle)
- text/rdf+nt (N-TRIPLES)

By default the URI of the content item being enhanced is a local, non de-referencable URI automatically built out of a hash digest of the binary content. Sometimes it might be helpful to



RESTful
is
Beautiful

localhost:8080/nuxeo/nxpath/default/default-domain/workspaces/News articles.1319058853893/Libya Rebel Muammar Gaddafi Search Advanced search

nuxeo • DM Home Document Management Admin Center Studio Administrator | SeamReload Search

default-domain Entities Sections Templates Workspaces AFP Arabic Another Workspace News articles

WorkList Clipboard No document in clipboard

Libya Rebels edge closer to Tripoli.html

Summary Edit Files Publish Relations Workflow Alerts Comments History Manage

Libyan rebels edged closer to the capital city of Tripoli on Sunday to help fellow mutineers inside the city who declared a final clash with leader Muammar Gaddafi. Following a night marred with gunfire, the rebels said that they controlled a handful of Tripoli's localities. With the rebels within about 25 km of Tripoli, Gaddafi's hold on power looks fragile. He labelled the rebels, who had been fighting for the past six months, as "rats" and said that he would not yield to their demands.

A coordinated revolt that rebels instantly after Muslim clerics called for advancing to the capital's peripol uprising, which is in its sixth-month. "The rebels may have risen too early," Oliver Miles, a former British ambassador said. "But quite the extent they think it has to happen." The rebels' advance toward the capital has been slow, with them and the capital. Government forces have taken over a burned-out tank, and some took refuge in buildings. Some rebels filled some walls with graffiti, or painted anti-Gaddafi slogans. In Benghazi, the rebels' main stronghold, everything was going according to plan. "There were no surprises, and others are coming in from outside," Ghoga, vice chairman of the rebel military committee, told reporters. Gaddafi — in hiding since the November 19 revolution — had said late yesterday that he had no intention of leaving Libya.

Moussa Ibrahim, in a briefing for journalists, said that the rebels' armed units defending Tripoli from the rebels wholeheartedly believe that if this city is captured, the blood will run everywhere; so they may as well fight to the end."



Muammar Muhammad al-Gaddafi (born 7 June 1942) is a Libyan revolutionary and the country's head of state from 1969 to the present day. Gaddafi became head of state after removing King Idris from power in a 1969 bloodless coup, after which he established the Libyan Arab Republic. His almost 42 years in power make him one of the longest-serving non-royal rulers in history.

Link to another entity

Mentions in current document

(...) the capital city of Tripoli on Sunday to help fellow mutineers inside the city who declared a final clash with leader **Muammar Gaddafi**. Following a night marred with gunfire, the rebels said that they controlled a handful of Tripoli's localities. With the rebels (...)

State Project Version 0.0

This document is **unlocked** | Lock

People

- Abdul Hafiz Ghoga
- Moussa Ibrahim
- Muammar Gaddafi
- Oliver Miles

Organizations

- Libya Rebels
- National Transition Council



localhost:8080/nuxeo/nxpath/default/default-domain/entities/Muammar Gaddafi@view_documents?tabIds=%20

nuxeo • DM Home Document Management Admin Center Studio Administrator | SeamReload Advanced search

default-domain Entities Sections Templates Workspaces

WorkList Clipboard No document in clipboard

Occurrences in documents

Title	Modified	Author	State
Libya Rebels edge closer to Tripoli.html	10/20/2011 12:40 AM	Administrator	Project
afp.com-20110615T020009Z-TX-PAR-QFU53.xml	10/20/2011 12:24 AM	Administrator	Project
afp.com-20110615T001023Z-TX-PAR-QFS22.xml	10/20/2011 12:24 AM	Administrator	Project
afp.com-20110615T071432Z-TX-PAR-QGG11.xml	10/20/2011 12:24 AM	Administrator	Project
afp.com-20110615T052303Z-TX-PAR-QGA07.xml	10/20/2011 12:24 AM	Administrator	Project
afp.com-20110615T044702Z-TX-PAR-QFY87.xml	10/20/2011 12:24 AM	Administrator	Project

Add link to new document

This Nuxeo DM instance is not registered to Nuxeo Connect: you won't be able to get the most recent fixes and stay up to date.
Nuxeo DM version: 5.4.3-SNAPSHOT. Register and enable Nuxeo Connect to benefit from automatic maintenance.

Muammar Gaddafi

Summary Edit Relations Alerts Comments History Manage

 Muammar Muhammad al-Gaddafi (born 7 June 1942) is a Libyan revolutionary and the country's head of state from 1969 to the present day. Gaddafi became head of state after removing King Idris from power in a 1969 bloodless coup, after which he established the Libyan Arab Republic. His almost 42 years in power make him one of the longest-serving non-royal rulers in history.

Also known as ムアンマル・アル=カッザーフィー, Muammar Gaddafi, مصطفى القذافي, Muammer Kaddafi, Muammar al-Gaddafi, Mu'ammar Gheddafi, Muamar el Gadafi, מוחמד א-גדהffi, 'ముఅమర్ కాధాఫి, 橋阿迈尔·卡扎菲, Каддафи, Муаммар

Entity Type Person

Remote knowledge base Muammar Gaddafi

> How to enable Nuxeo Connect
> Subscribe to Nuxeo Connect





Stanbol

=

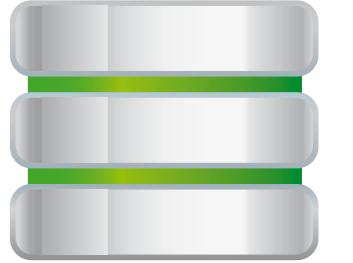
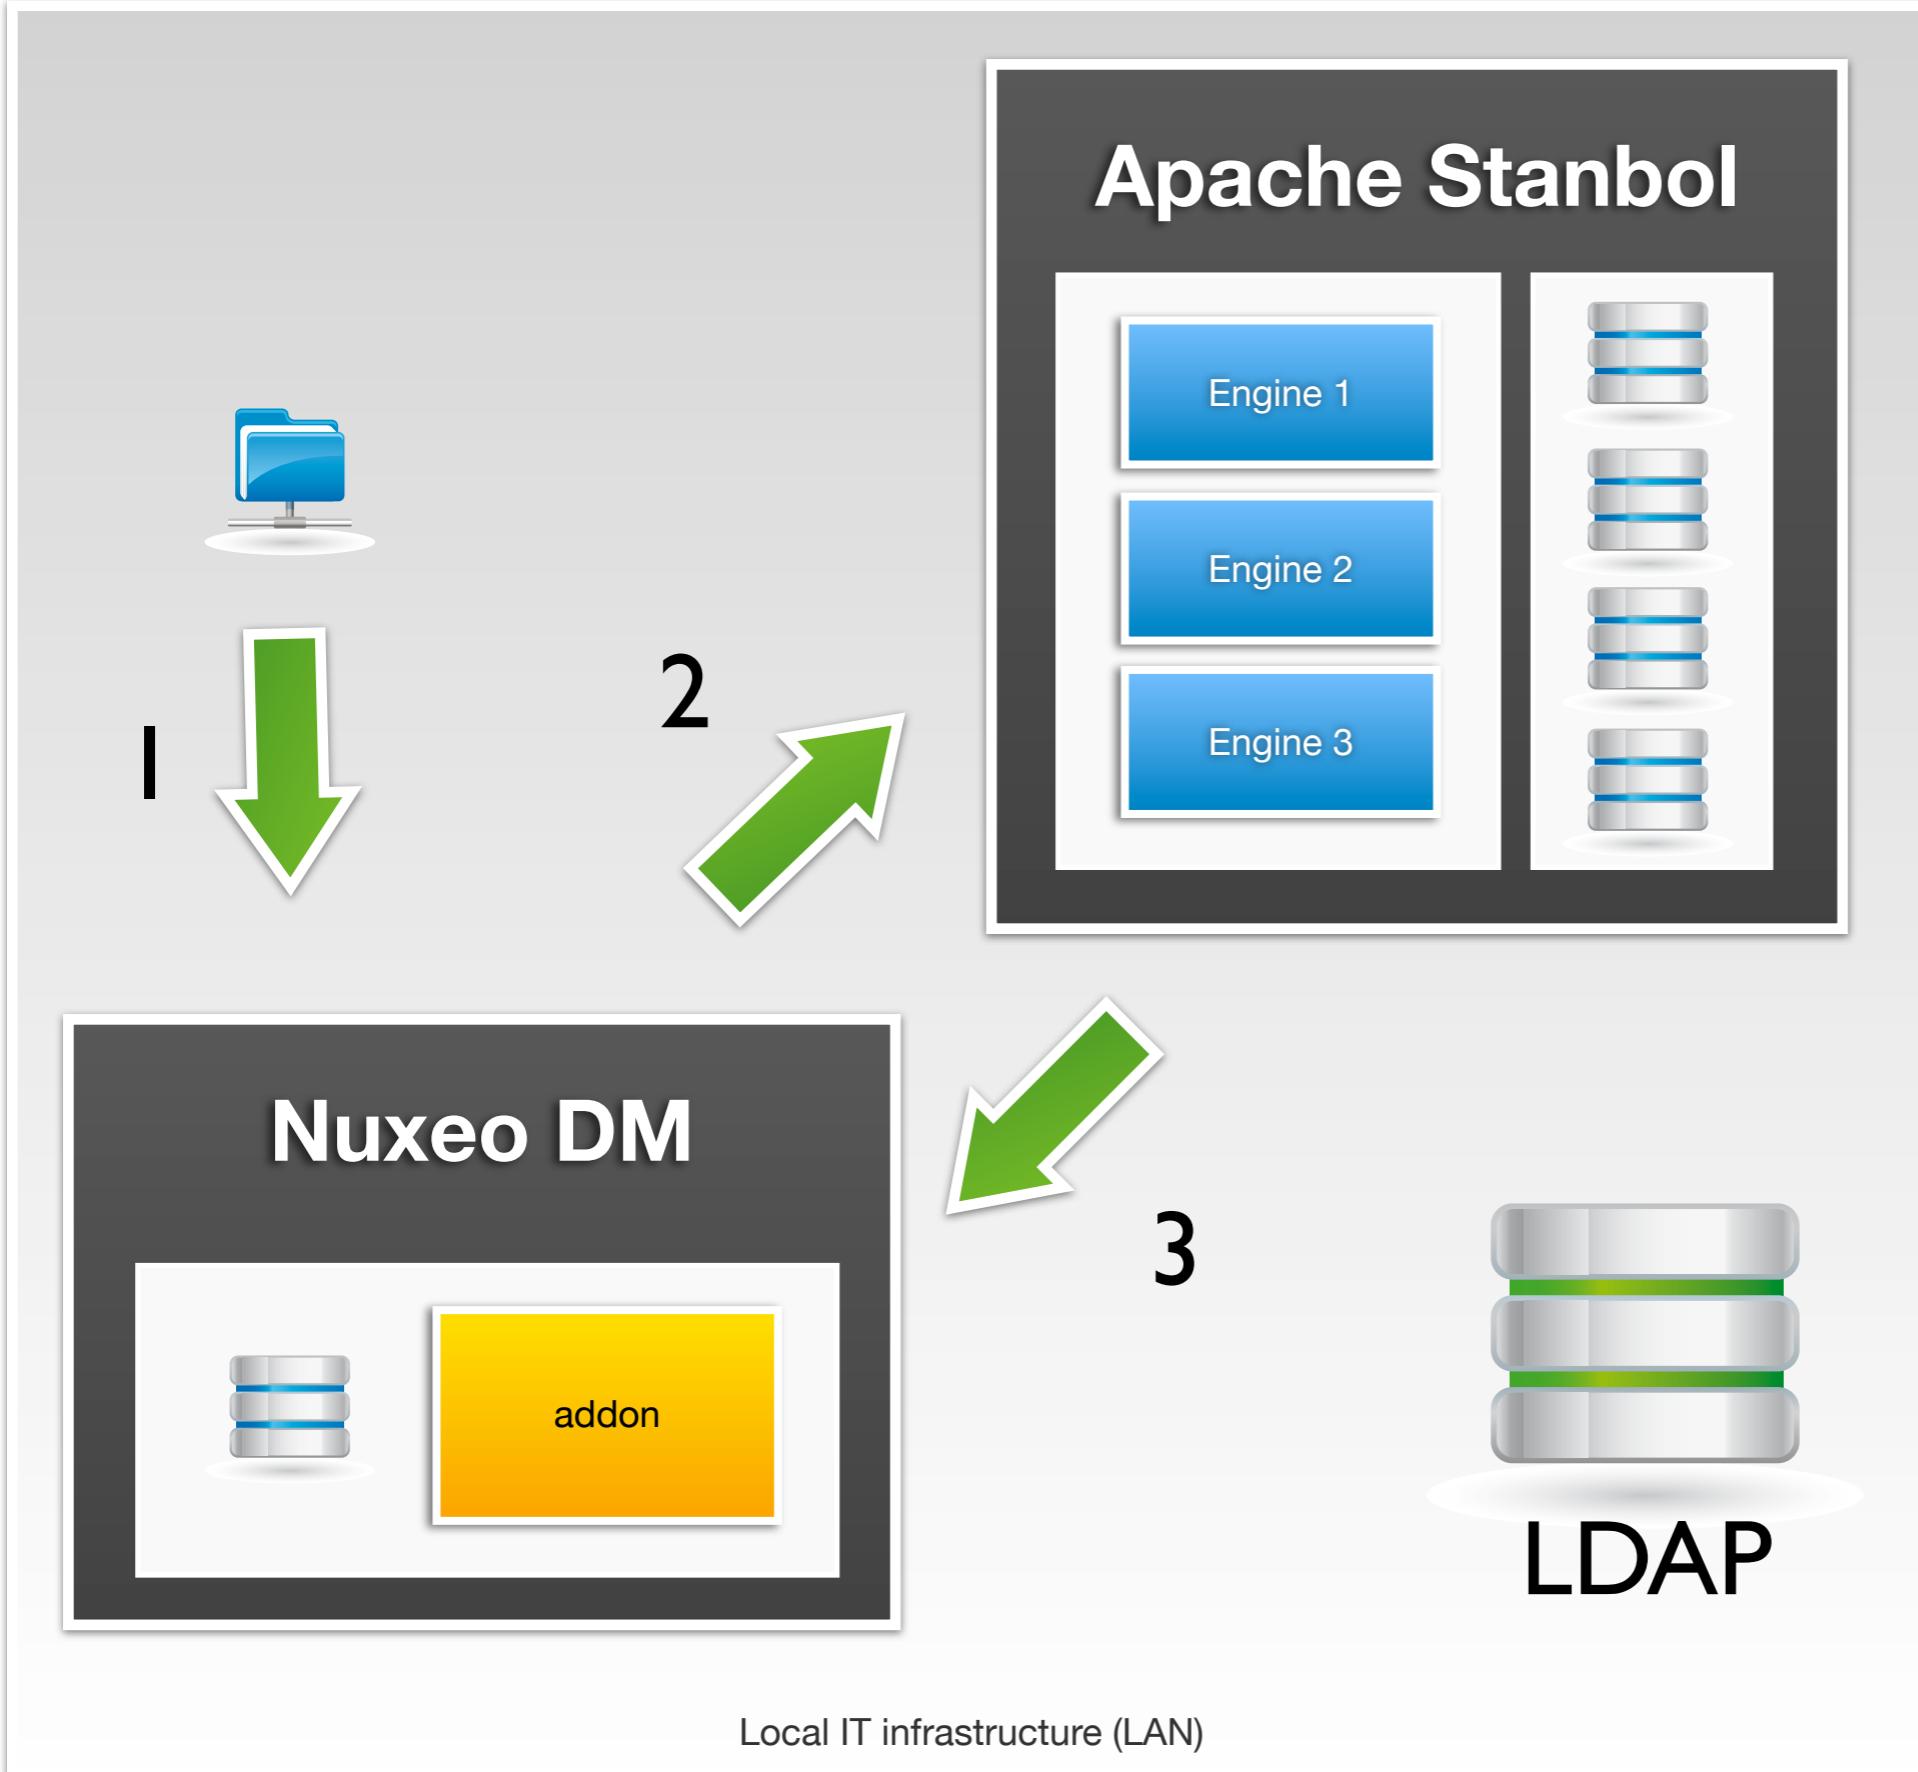
Semantic Engines
(Apache OpenNLP)

+

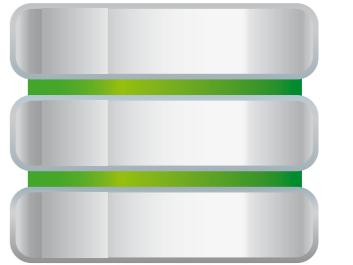
Fast Linked Data local index
(Apache Solr)

+

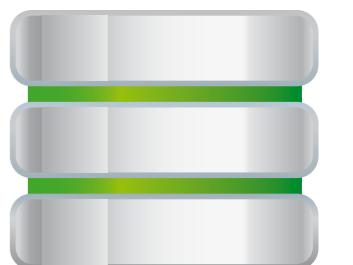
Semantic Rule Engine
(Apache Jena)



DBpedia



Freebase



Geonames

How to build engines?

Training statistical models for NER with Wikipedia and DBpedia

- Extract **sentences with link positions** in Wikipedia articles
- DBpedia to the **find type of the target entity** (Person, Location, Organization)
- **Apache Pig scripts** to compute the **join + format** the result as training files for OpenNLP
- **Apache OpenNLP** to build and evaluate the models
- **Apache Hadoop** for distributed processing
- **Apache Whirr** for deployment and management on Amazon EC2 cluster



```
-- Register the project jar to use the custom loaders and UDFs
REGISTER $PIGNLPROC_JAR

parsed = LOAD '$INPUT'
    USING pignlproc.storage.ParsingWikipediaLoader('$LANG')
    AS (title, wikiuri, text, redirect, links, headers, paragraphs);

-- filter and project as early as possible
noredirect = FILTER parsed by redirect IS NULL;
projected = FOREACH noredirect GENERATE title, text, links, paragraphs;

-- Extract the sentence contexts of the links respecting the paragraph
-- boundaries
sentences = FOREACH projected
    GENERATE title, flatten(pignlproc.evaluation.SentencesWithLink(
        text, links, paragraphs));

stored = FOREACH sentences
    GENERATE title, sentenceOrder, linkTarget, linkBegin, linkEnd, sentence;

-- Ensure ordering for fast merge with type info later
ordered = ORDER stored BY linkTarget ASC, title ASC, sentenceOrder ASC;
STORE ordered INTO '$OUTPUT/$LANG/sentences_with_links';
```



```
-- Load wikipedia, instance types and redirects from DBpedia dumps
wikipedia_links = LOAD '$INPUT/wikipedia_links_${LANG}.nt'
USING pignlproc.storage.UriUriNTriplesLoader(
  'http://xmlns.com/foaf/0.1/primaryTopic')
AS (wikiuri: chararray, dburi: chararray);

wikipedia_links2 = FILTER wikipedia_links BY wikiuri IS NOT NULL;

-- Load DBpedia type data and filter out the overly generic owl:Thing type
instance_types =
  LOAD '$INPUT/instance_types_en.nt'
  USING pignlproc.storage.UriUriNTriplesLoader(
    'http://www.w3.org/1999/02/22-rdf-syntax-ns#type')
  AS (dburi: chararray, type: chararray);

instance_types_no_thing = FILTER instance_types BY type NEQ 'http://www.w3.org/2002/07/owl#Thing'
joined = JOIN instance_types_no_thing BY dburi, wikipedia_links2 BY dburi;
projected = FOREACH joined GENERATE wikiuri, type;

-- Ensure ordering for fast merge with sentence links
ordered = ORDER projected BY wikiuri ASC, type ASC;
STORE ordered INTO '$OUTPUT/${LANG}/wikiuri_to_types';
```



```
sentences = LOAD '$INPUT/$LANG/sentences_with_links'
AS (title: chararray, sentenceOrder: int, linkTarget: chararray,
linkBegin: int, linkEnd: int, sentence: chararray);

wikiuri_types = LOAD '$INPUT/$LANG/wikiuri_to_types'
AS (wikiuri: chararray, typeuri: chararray);

-- load the type mapping from DBpedia type URI to OpenNLP type name
type_names = LOAD '$TYPE_NAMES' AS (typeuri: chararray, typename: chararray);

-- Perform successive joins to find the OpenNLP typename of the linkTarget
joined = JOIN wikiuri_types BY typeuri, type_names BY typeuri USING 'replicated';
joined_projected = FOREACH joined GENERATE wikiuri, typename;
joined2 = JOIN joined_projected BY wikiuri, sentences BY linkTarget;

result = FOREACH joined2
GENERATE title, sentenceOrder, typename, linkBegin, linkEnd, sentence;

-- Reorder and group by article title and sentence order
ordered = ORDER result BY title ASC, sentenceOrder ASC;
grouped = GROUP ordered BY (title, sentenceOrder);

-- Convert to the OpenNLP training format
opennlp_corpus =
FOREACH grouped
GENERATE opennlp_merge(
    ordered.sentence, ordered.linkBegin, ordered.linkEnd, ordered.typename);
```

```
$ opennlp TokenNameFinderEvaluator -encoding utf-8 \
  -model fr-ner-location \
  -data ~/data/fr/opennlp_location/test
```

Performance evaluation for NER on a French extraction with 100k sentences

class	precision	recall	f1-score
location	0.87	0.74	0.80
person	0.80	0.68	0.74
organization	0.80	0.65	0.72

Performance evaluation for NER on a English extraction with 100k sentences

class	precision	recall	f1-score
location	0.77	0.67	0.71
person	0.80	0.70	0.75
organization	0.79	0.64	0.70



Training statistical models for topic classification from Wikipedia and DBpedia

- Filter category tree from **DBpedia SKOS entries** (~500k)
- **Pig scripts** to compute the **joins with articles abstracts** for all the articles categorized in Wikipedia
- Export as 2.8GB TSV file to be indexed in **Apache Solr**
- Use Solr **MoreLikeThisHandler** to find the top 3 most related Wikipedia category for any kind of text
- Apache Whirr & Hadoop for deployment and management on Amazon EC2 cluster



Wrap Up on Recent Work

- Full offline mode: Stanbol EntityHub
- Multi-lingual Indexes
- New UI for occurrences reviews
- Temis Luxid Annotation Factory integration



What's next?

- Stanbol and Temis connection in Admin Center
- Embedded Stanbol mode for easy deployment
- More OpenNLP models for more languages
- Finalize topic classification - handle hierarchy
- Tight integration with Nuxeo DM search features



Thank you for your attention!



48

NUXEO WORLD PARIS 2011

